

Hedgehog Installation Guide 2.0.0b1

- 1. Installation
 - 1.1. Requirements
 - 1.1.1. Platform support
 - 1.1.2. Database support
 - 1.2. Hedgehog dependencies
 - 1.2.1. Required packages
 - 1.2.2. Additional R packages
 - 1.3. Install Hedgehog
 - 1.3.1. Install and build code
 - 1.3.2. System users and groups
 - 1.3.3. Configure Hedgehog
 - 1.3.4. Configure apache
 - 1.3.5. Create and populate the database
 - 1.3.6. Test the Hedgehog front end
- 2. Running Hedgehog
 - 2.1. Import data
 - 2.1.1. Importing historical .DAT data
 - 2.1.2. Importing real-time .XML data
 - 2.1.2.1. Manually
 - 2.1.2.2. Automatically
 - 2.2. Maintenance
 - 2.2.1. Cron jobs

The following instructions describe how to install Hedgehog 2.0.0b1 on Ubuntu 14.04 Server and set up imports of data.



The server and PostgreSQL must both be configured to use UTC!

1. Installation

1.1. Requirements

1.1.1. Platform support

Version 2.0.0b1 of Hedgehog is currently only supported on Ubuntu 14.04 LTS Server.

1.1.2. Database support

Hedgehog has been tested with PostgreSQL 9.3.X

1.2. Hedgehog dependencies

1.2.1. Required packages

```
sudo add-apt-repository ppa:opencpu/rapache
sudo apt-get update
sudo apt-get install r-base r-base-core r-base-dev libcairo2-dev libxt-dev \
    r-cran-ggplot2 r-cran-dbi r-cran-cairodevice r-cran-reshape \
    r-cran-digest postgresql postgresql-client \
    postgresql-server-dev-9.3 postgresql-contrib pgbouncer \
    libboost-all-dev libtool libpqxx-dev apache2-mpm-prefork \
    apache2-prefork-dev git libapache2-mod-r-base
```

1.2.2. Additional R packages

Some R packages are needed that are not available through apt, they can be installed via R.

You will be prompted to choose a mirror site, or you can specify a repo by using a command of the form

```
install.packages("name", repos='http://cran.rstudio.com/')
```

Install:

```
sudo R
install.packages(c("brew", "Cairo", "googleVis", "RPostgreSQL", "R.utils", "yaml"))
q()
```

If you are prompted to *save workspace image y/n/c*, choose *no*.

The package versions that Hedgehog has been tested against are listed below alongside each package.

R Package	Supported Version
brew	1.0-6
Cairo	1.5-6
googleVis	0.5.5
RPostgreSQL	0.4
R.utils	1.33.0
yaml	2.1.13

1.3. Install Hedgehog

1.3.1. Install and build code

Download the 2.0.0b1 release from github: <https://github.com/dns-stats/hedgehog/archive/2.0.0b1.tar.gz>

```
tar -xzf hedgehog-2.0.0b1.tar.gz
cd hedgehog-2.0.0b1
./autogen.sh
mkdir build
cd build
../configure
make
sudo make install
```

If no <prefix> is specified in *configure* then it defaults to */usr/local/*

- You may notice the following message during the installation step: *"Error in find.package(pkgs, lib) : there is no package called 'RPostgreSQLHelper' Execution halted"*. This can be safely ignored.



For large installations it may be preferable to store the data files on a separate partition to the database. See the *User guide for an overview of the Hedgehog directory structure.

In this case you may wish to mount a disk here or make use of a symbolic link at this stage of the installation.

1.3.2. System users and groups

Two system users are required for Hedgehog:

- Database owner - this is the user that will own the database created by hedgehog and the top level datafile directories (default user in this example is 'hedgehog')
- Read user - this should be the apache user which requires only read access to the database (default apache user in this example is called 'www-data')

```
# create a system user and group called hedgehog
sudo addgroup --system hedgehog
sudo adduser --system --ingroup hedgehog hedgehog
# put www-data user in group hedgehog
sudo adduser www-data hedgehog
# put hedgehog user in group www-data
sudo adduser hedgehog www-data
```

1.3.3. Configure Hedgehog

- Edit the `<prefix>/etc/hedgehog/hedgehog.yaml` configuration file to make sure the user names match the users specified in the previous step.
- Also configure the database parameters (port, name) as required. (Note a limitation of 2.0 is that the database and web front end must run on the same machine)

```
database:
  port      : 5432      # specify port for the postgresql DB.
  name      : hedgehog  # specify dbname for the postgresql DB.
  owner     : hedgehog  # specify an owner for the postgresql DB.
                  # this user will also run refile_and_grok.
  read_user : www-data  # specify a read user for the postgresql DB.
                  # this is typically the apache user.
```



Do not edit the 'directories' section of this file as in 2.0 it is auto-generated.

Note that there is also a sample conf file in the same directory to retrieve default settings: `<prefix>/etc/hedgehog/hedgehog.yaml.sample`.

- Change the ownership of the directory specified below so it is owned by the same user as the database, so for example in a default install:

```
sudo chown -R hedgehog:hedgehog <prefix>/var/hedgehog/
```

- Also check the parameters in the `<prefix>/etc/hedgehog/hedgehog_gui.yaml` file, which specifies parameters controlling the behaviour of the web front end. See the "Plot Caching" section in the user guide for a more detailed description of when plots are cached.

```

---
# YAML config for hedgehog GUI.
# NOTE: If this file is changed then apache must be restarted for the changes to take
effect
www:
  default_plot_type          : interactive # 'static'      -> png plots
                                # 'interactive' -> googlevis plots
  default_interactive_plot_type : flash    # 'flash' -> plot requires flash
                                # 'svg'      -> plot is SVG/VML and does
                                # not require flash (but with svg plots
                                # some legends do not wrap properly)
  use_plot_caching           : 1          # '1' -> true, use cached plots when
                                # possible
                                # '0' -> false, never use cached plots
  caching_delay_in_hours     : 1          # If 'use_plot_caching=1' then only
                                # plots with an end time earlier than
                                # this number of hours ago are cached.
                                # More recent plots are not cached as
                                # data may still be being imported
  presentation_delay_in_hours : 0          # Number of hours behind now for which
                                # the GUI will display data

```

1.3.4. Configure apache

You will probably need to edit the `/etc/apache2/apache2.conf` file to enable access to the install directory by adding `<Directory <prefix>/share/hedgehog>` and `<Directory <prefix>/var/hedgehog/www>` elements.

Depending on your configuration you may want to change the name of the `hedgehog.conf` file or disable the default site using the following command:

```
sudo a2dissite 000-default.conf
```

- Add the Hedgehog configuration files to apache and enable the site:

```
sudo cp <prefix>/share/hedgehog/conf/hedgehog.conf /etc/apache2/sites-available/
sudo a2ensite hedgehog.conf
```

- Alter the permissions (add the line 'umask 002' to the apache envvars file):

```
> sudo vi /etc/apache2/envvars
# Enable extra permissions for Hedgehog
umask 002
```



Uploading XML via webdav using client certificates

For installs that **want to upload XML via webdav using client certificates** then:

```
sudo vi /etc/apache2/conf-available/hedgehog.conf
```

and uncomment the "*Alias /data*" line and following directory clause in this file.

Alter the apache umask so that www-data group members (i.e. hedgehog) can process the xml files :

Clearly you will also need to issue certificates to both Apache and to every collector. Help with this can be found User Guide.



apache/rapache write some of their logs to user.* so it can be useful to change the syslog config:

```
sudo vi /etc/rsyslog.d/50-default.conf
```

apache/rapache write some of their logs to user.* so it can be useful to change the syslog config:

Uncomment the line beginning 'user.*'.

- Finally, restart apache:

```
sudo service apache2 restart
```

1.3.5. Create and populate the database



For this version of Hedgehog the servers and nodes to be processed and displayed must be specified manually.

- Edit the the `<prefix>/etc/hedgehog/nodes.csv` to specify the servers, nodes and grouping to be used (example format is provided with entries commented out).
 - Note that the current GUI layout is optimised for nodes with short names (<6 characters) of the same length
- Then run the command below noting the following:
 - If you have historic data to import then use the `-m` flag to specify the month of the oldest data that will need importing. Otherwise the database tables will be created to hold data from this month onwards.
 - This script assumes there is a default postgres user called 'postgres' (with no password configured) that has sufficient privileges to create databases and users.
 - If the the default postgres user is not called 'postgres' then use the `-u` parameter to specify the name of the database user to use to create the Hedgehog database
 - If the postgres user requires a password, use the `-p` flag to be prompted to enter the password.
 - Note that this script will also create the directory structure for all the specified servers and nodes under the data directory if it does not exist

```
sudo <prefix>/sbin/hedgehog_database_create.sh
```

1.3.6. Test the Hedgehog front end

At this point you should be able to see the servers and nodes in the web front end at the URL `http://<server-name>/hedgehog`

2. Running Hedgehog

2.1. Import data

Hedgehog can process data in the following 3 ways:

Source format	Output format	
---------------	---------------	--

XML	Database	For real time uploads
DAT	Database	For import of historic data
XML	DAT	For backwards compatibility with DSC

In each case the `<prefix>/bin/refile_and_grok.sh` script is used, it is simply given different parameters:

```
> refile_and_grok.sh -h

refile_and_grok - finds all input files in the working directory and processes to
output format

-w Working directory to search for input files (default: <prefix>/var/hedgehog/data)
-i Input file format <XML|DAT> (default: XML)
-o Output file format <DAT|DB> (default: DB)
-c Non-interactive mode - use this flag when being run by a cron job
-s Start date from which to process incoming data (XML input only)
-r Disable processing of rssac datasets. Default is to process all datasets.
-R Reserved processors. Number of CPUS processors to exclude from import (default 0).
-h Show this help.
```

2.1.1. Importing historical .DAT data

```
sudo -u hedgehog <prefix>/bin/refile_and_grok.sh -i DAT
```

Be aware that this can take a long time if there is a significant amount of historic data and it may be advisable to run this in stages.

2.1.2. Importing real-time .XML data

2.1.2.1. Manually

- This can be done manually by running the `refile_and_grok.sh` script (consider running this `nohup` as it may take a while depending on how much data there is to process).

```
sudo -u hedgehog <prefix>/bin/refile_and_grok.sh
```

- A snapshot of the progress of the data import can be generated by running the command below:

```
sudo -u hedgehog <prefix>/bin/hedgehog_import_create_summary.sh
```

2.1.2.2. Automatically

- Configure a regular cron job for `refile_and_grok.sh` as shown below

2.2. Maintenance

2.2.1. Cron jobs

In 2.0 several cron jobs need to be configured. They should run as the 'hedgehog' system user.

```
sudo -u hedgehog crontab -e
```

Below is an example crontab for a typical system.



Note that the *hedgehog_manage_partitions_cron.sh* script **MUST** be configured to run at least once a month.

```
# REQUIRED:
# Import XML data every 15 mins
00,15,30,45 * * * * <prefix>/bin/refile_and_grok.sh -c >>
/home/hedgehog/refile_and_grok_xml_to_db.sh.log 2>&1
# Twice monthly job to make sure the DB tables for next month are created
# ahead of time
0 6 15,28 * * <prefix>/bin/hedgehog_manage_partitions.sh >>
/home/hedgehog/hedgehog_manage_partitions.sh.log 2>&1

# OPTIONAL:
# Daily job to create cached plots for the previous day to make loading common plots
# quicker. Run a few hours after midnight so all data is uploaded.
0 4 * * * <prefix>/bin/hedgehog_plotcache_generate_cached_plots.sh -D >>
/home/hedgehog/hedgehog_plotcache_generate_cached_plots.sh.log -D 2>&1
# Daily job to generate RSSAC reports. By default report is generated
# for a single day 1 week ago
0 1 * * * <prefix>/bin/hedgehog_rssac_generate_reports.sh -D >>
/home/hedgehog/hedgehog_rssac_generate_reports.sh.log 2>&1
# Monthly job to tar up processed xml directories
0 2 1 * * <prefix>/bin/hedgehog_datafiles_rm_empty_xml_dirs.sh -D >>
/home/hedgehog/hedgehog_datafiles_rm_empty_xml_dirs.sh.log 2>&1
# Monthly job to remove empty xml directories that are older than 7 days old
0 2 7 * * <prefix>/bin/hedgehog_datafiles_tar_old_xml.sh -D >>
/home/hedgehog/hedgehog_datafiles_tar_old_xml.sh.log 2>&1
```