

Hedgehog Installation from Ubuntu Packages Guide 2.1.0rc1

- 1. Requirements
 - 1.1. Platform support
 - 1.2. Database support
 - 1.3. General
 - 1.3.1. Home directory permissions
- 2. Hedgehog Installation
 - 2.1. Hedgehog packages
 - 2.2. Web GUI only - Manual install of R packages
- 3. Hedgehog configuration
 - 3.1. Database configuration
 - 3.1.1. Create a database
 - 3.2. Data Manager configuration
 - 3.2.1. Specify the nodes/servers
 - 3.2.2. Directory permissions
 - 3.2.3. Create the database tables
 - 3.2.4. (Optional) Webdav upload
 - 3.3. Web GUI configuration
 - 3.3.1. Apache configuration
- 4. Importing data
 - 4.1. Importing historical .DAT data
 - 4.2. Importing real-time .XML data
 - 4.2.1. Manually
 - 4.2.2. Automatically
 - 4.3. Importing zone-size and load-time data for RSSAC
- 5. Cron jobs
 - 5.1. Data Manager component
 - 5.2. Web GUI component

The following instructions describe how to install Hedgehog 2.1.0rc1 from Ubuntu Packages on Ubuntu 14.04 Server and set up imports of data.

1. Requirements

1.1. Platform support

Version 2.1.0rc1 of Hedgehog is currently only supported on Ubuntu 14.04.3 LTS Server.

1.2. Database support

Hedgehog has been tested with PostgreSQL 9.3.X. The installation of this of this is not covered in detail since it may or may not be co-located with the other hedgehog components. If not already installed, it will be installed as a dependency of the hedgehog-database package.

The Hedgehog servers and PostgreSQL must both be configured to use UTC!

1.3. General

1.3.1. Home directory permissions

Some users set the permissions on their home directory such that other users cannot read or execute that directory. In this case when running scripts that update the database as the *hedgehog* or *postgres* (user via 'sudo -u <user>') meaningless messages are generated from psql about being unable to cd into logged in users home directory. To avoid this spurious output change directory to one that allows other users to read/execute (e.g. /tmp) and run the scripts from there.

2. Hedgehog Installation

2.1. Hedgehog packages

The Hedgehog PPA is located at

```
sudo add-apt-repository ppa:dns-stats/presenter
```

Hedgehog comprises 3 main components which can be run on the same or different machines. Configure the additional PPA's as required and then install the main packages from the list below:

Component	Description	Main Package	Optional Package	Additional Required PPAs
Database	PostgreSQL database for Hedgehog	hedgehog-database		sudo add-apt-repository ppa:ondrej/pkg-nlnetlabs
Data Manager	Scripts for database population and management	hedgehog-data-manager		
	Scripts for XML/DAT processing			
	(Optional) Apache config for webdav based XML uploads		hedgehog-webdav-upload	
Web GUI	Web GUI front-end	hedgehog-gui		sudo add-apt-repository ppa:opencpu/rapache

Notes:

- Multiple Web GUI's can be connected to the same database.
- The upload mechanism for XML/DAT files is for the user to decide (ssh, rsync, Apache, etc.). A package for upload using Apache is provided for convenience.

For reference the other Hedgehog packages are:

Description	Required by	Package
Common package	All main hedgehog packages	hedgehog-common
RPostgresHelper	Web Gui	r-hedgehog-rpostgresqlhelper
Documentation		hedgehog-doc

Two system users are used by Hedgehog, which are created by the packages:

User	Default	Component	Note	Created by
Database owner	hedgehog	Data Manager	This is the user that will own the database created by Hedgehog and the top level datafile directories.	hedgehog-common package
Database read user	www-data	Web front-end	This defaults to the same as the default apache user	apache2 package installation

2.2. Web GUI only - Manual install of R packages

These instructions are for R packages that must be installed using R's built-in compilation tools (there is no Ubuntu package for them).

To install interactively: (Some can repositories don't contain packages for all versions of R so if this fails quit R and re-try a different repo)

```
sudo R
install.packages(c("brew", "Cairo", "googleVis", "R.utils", "yaml", "plyr"))
q()
# If you are prompted to save workspace image y/n/c, choose no.
```

For scripting purposes a repo can be specified by using a command of the form:

```
install.packages("name", repos='http://cran.rstudio.com/')
```

Hedgehog is tested against version 3.1.1 of R. The package versions that Hedgehog has been tested against are listed below alongside each package.

R Package	Supported Version
brew	1.0-6
Cairo	1.5-9
googleVis	0.5.10
R.utils	2.2.0
yaml	2.1.13
plyr	1.8.3

3. Hedgehog configuration

- For the Web GUI and Data Manager components, ensure the `/etc/hedgehog/hedgehog.yaml` configuration file contains usernames and passwords that will match those configured in the database (see next section).
- Also configure the database parameters (host, port and name) as required.

```
database:
  host      : /var/run/postgresql # specify a host for the postgresql
DB. If                                           # this begins with a slash, it
                                                # directory in which the socket
specifies the                                # stored.
file is
  port      : 5432      # specify port for the postgresql DB.
  name      : hedgehog  # specify dbname for the postgresql DB.
  owner     : hedgehog  # specify a user to own the postgresql DB.
                        # [Required for Data Manager component]
  owner_pass :          # specify a password for the owner user if
needed.
  read_user  : www-data # specify a read user for the postgresql DB.
                        # [Required for Web front-end component]
  read_pass  :          # specify a password for the read user if
needed.
```

Do not edit the 'directories' section of this file as it is auto-generated.

Depending on your PostgreSQL configuration you may need to add passwords to this file. If you do then be aware that, by default, this file is readable by all users. You should restrict access to just the hedgehog user on a Data Manager only machine or the hedgehog and www-data user on a combined Data Manager and Web GUI.

3.1. Database configuration

3.1.1. Create a database

Ask your DBA to create the necessary database. This is a script to help them. It creates the DATABASE, SCHEMA, FUNCTIONS, EXTENSIONS, LANGUAGES, USERS and ROLES needed to run hedgehog (using default values), and can optionally accept a user specified database name and read/write user names.

```
sudo -u postgres /usr/bin/hedgehog_database_create

#If you want to use passwords do something like this or use peer
authentication
sudo -u postgres psql
alter user hedgehog password 'hedgehog123';
alter user "www-data" password 'www123';
```

Modify the PostgreSQL configuration as so:

```
sudo vi /etc/postgresql/9.3/main/postgresql.conf
# uncomment and set 'extra_float_digits = 1'
sudo service postgresql reload
```

We recommend that the 'pgtune' tool is used to obtain an optimal configuration for PostgreSQL. For example:

```
pgtune -i /etc/postgresql/9.3/main/postgresql.conf -c 200 -T DW
```

Also some queries to the database trigger DNS lookups from functions in the database. It may be optimal to run a local resolver such as Unbound.

3.2. Data Manager configuration

3.2.1. Specify the nodes/servers

For this version of Hedgehog the servers and nodes to be processed and displayed must be specified manually as described here.

- Edit the the `/etc/hedgehog/nodes.csv` file to specify the servers, nodes and grouping to be used (example format is provided with entries commented out).
- Note that the current GUI layout is optimised for nodes with short names (<6 characters) of the same length

3.2.2. Directory permissions

The `/var/lib/hedgehog/data/` is used to store incoming XML files

You may also need to alter the permissions on this directory to allow uploads via your chosen mechanism

3.2.3. Create the database tables

Run the command below noting the following:

- If you have historic data to import then use the `-m` flag to specify the month of the oldest data that will need importing. Otherwise the database tables will be created to hold data from this month onwards.
- Note that this script will also create the directory structure for all the specified servers and nodes under the `data` directory if it does not exist
- (Note the insertion of the GeoIP data can take some time)

```
sudo -u hedgehog /usr/bin/hedgehogctl database_init
```

Then run the following script to make sure the list of delegated TLDs in the database is up to date.

```
sudo -u hedgehog /usr/bin/hedgehogctl database_update_tlds_from_zone  
2>/dev/null
```

3.2.4. (Optional) Webdav upload

If using webdav to upload XML files then add the following to the `/etc/apache2/envvars` file:

```
umask 002
```

Then enable the dependencies:

```
sudo a2enmod dav  
sudo a2enmod ssl  
suod a2enmod dav_fs
```

And then enable the hedgehog webdav site:

```
sudo a2ensite hedgehog-webdav-upload
```

3.3. Web GUI configuration

Check the parameters in the `/etc/hedgehog/hedgehog_gui.yaml` file, which specifies parameters controlling the behaviour of the web front end. See the "Plot Caching" section in the user guide for a more detailed description of when plots are cached.

```

---
# YAML config for hedgehog GUI.
# NOTE: If this file is changed then apache must be restarted for the
changes to take effect
www:
  default_plot_type          : interactive # 'static'      -> png
plots
                                # 'interactive' ->
googlevis plots
  default_interactive_plot_type : svg          # 'flash' -> plot requires
flash
                                # 'svg'    -> plot is
SVG/VML and
                                # does not require flash
(but with svg
                                # plots some legends do
not wrap properly)
  default_node_grouping        : instance      # choose from 'none',
'instance', 'city' or
                                # 'country'
  use_plot_caching             : 1             # '1' -> true, use cached
plots when possible
                                # '0' -> false, never use
cached plots
  caching_delay_in_hours       : 1             # If 'use_plot_caching=1'
then only plots with
                                # an end time earlier than
this number of
                                # hours ago are cached.
More recent plots are
                                # not cached as data may
still be importing
  presentation_delay_in_hours  : 0             # Number of hours behind
now for which the
                                # GUI will display data
  support_url                  :                # configurable target for
"Support" external
                                # link on Homepage. The
default dns-stats.org
                                # issue tracker used if
left blank

```

3.3.1. Apache configuration

Depending on your exact installation choices and apache configuration you may need to disable the default site using the following command:

```
sudo a2dissite 000-default.conf
```

- Add the Hedgehog configuration files to apache and enable the site (this file name can be changed if required to match any local apache policy):

```
sudo a2ensite hedgehog.conf
```

apache/rapache write some of their logs to user.* so it can be useful to change the syslog config:

```
sudo vi /etc/rsyslog.d/50-default.conf
```

Uncomment the line beginning 'user.*'.

- Finally, reload apache:

```
sudo service apache2 reload
```

At this point you should test that you can see the servers and nodes in the web front end at the URL <http://<server-name>/hedgehog>

4. Importing data

Hedgehog can process data in the following 3 ways:

Source format	Output format	
XML	Database	For real time uploads
DAT	Database	For import of historic data
XML	DAT	For backwards compatibility with DSC

In each case the `/usr/bin/refile_and_grok` script is used, it is simply given different parameters:

```
> refile_and_grok -h
```

```
refile_and_grok - finds all input files in working directory and  
processes to output format
```

```
-w Working directory to search for input files (default: )  
-i Input file format <XML|DAT> (default: XML)  
-o Output file format <DAT|DB> (default: DB)  
-c Non-interactive mode - use this flag when being run by a cron job  
-s Start date from which to process incoming data (XML input only)  
-r Disable processing of rssac data. Default is to process all data.  
-R Reserved processors. Number of CPUS processors to exclude from import  
(default 0).  
-a Append output to the refile_and_grok.stdout file (default is  
overwrite)  
-h Show this help.
```

4.1. Importing historical .DAT data

```
sudo -u <DB_OWNER> /usr/bin/refile_and_grok -i DAT
```

Be aware that this can take a long time if there is a significant amount of historic data and it may be advisable to run this in stages.

4.2. Importing real-time .XML data

4.2.1. Manually

- This can be done manually by running the *refile_and_grok* script (consider running this *nohup* as it may take a while depending on how much data there is to process).

```
sudo -u <DB_OWNER> /usr/bin/refile_and_grok
```

- A snapshot of the progress of the data import can be generated by running the command below:

```
sudo -u <DB_OWNER> /usr/bin/hedgehogctl datafiles_create_summary
```

4.2.2. Automatically

- Configure a regular cron job for *refile_and_grok* as shown below

4.3. Importing zone-size and load-time data for RSSAC

To do this run the *rssacd* demon specifying the server of interest. This listens for NOTIFY messages and after receiving one:

- obtains the zone size from the notifying server by performing an XFR and
- probes each node listed for the configured server with an IP address to calculate the zone load time

This requires that management IP addresses are configured for the nodes via the *nodes.csv* file (also note that *rssacd* needs restarting if new nodes are added).

rssacd can be run from the command line for testing, or using an init script.

From the command line:

```
/usr/sbin/rssacd --log ~/var/log/hedgehog/rssacd.log -s <server-name> -z  
<fully_qualified_zone_name>
```

From an init script:

- Configure the server name, TSIG key, etc. in the */etc/hedgehog/rssac.conf* file

Then run:

```
update-rc.d rssacd defaults
```

5. Cron jobs

In 2.1.0 several cron jobs need to be configured.

5.1. Data Manager component

Below is an example crontab for a typical data manager install (`sudo -u <DB_OWNER> crontab -e`).

Note that the `database_manage_partitions` script MUST be run at least once a month to create the tables for next month or the import will fail.

```
# REQUIRED:
# Import XML data every 15 mins
00,15,30,45 * * * * /usr/bin/refile_and_grok -c >>
/var/log/hedgehog/refile_and_grok_xml_to_db.log 2>&1
# Twice monthly job to make sure the DB tables for next month are
created
# ahead of time
0 6 15,28 * * /usr/bin/hedgehogctl database_manage_partitions >>
/var/log/hedgehog/database_manage_partitions.log 2>&1

# OPTIONAL:
# Daily jobs to process RSSAC data. By default data is processed
# for a single day 1 week ago. Must be run before the
rssac_generate_reports script (see Web GUI Component cron jobs)
0 1 * * * /usr/bin/hedgehogctl database_process_rssac_data -D >>
/var/log/hedgehog/database_process_rssac_data.log 2>&1
# Weekly job to update the delegated TLDs from the IANA database.
0 2 * * 0 <prefix>/bin/hedgehogctl database_update_tlds_from_zone >>
/var/log/database_update_tlds_from_zone.log 2>&1
# Monthly job to update the geoIP database.
0 3 1 * * <prefix>/bin/hedgehogctl database_update_geoip >>
/var/log/database_update_geoip.log 2>&1
# Monthly job to remove empty xml directories that are older than 7 days
old
0 2 1 * * /usr/bin/hedgehogctl datafiles_rm_empty_xml_dirs -D >>
/var/log/hedgehog/datafiles_rm_empty_xml_dirs.log 2>&1
# Monthly job to tar up processed xml directories
0 2 7 * * /usr/bin/hedgehogctl datafiles_tar_old_xml -D >>
/var/log/hedgehog/datafiles_tar_old_xml.log 2>&1
```

5.2. Web GUI component

Below is an example crontab for a typical web front-end install (`sudo -u <DB_READ_USER> crontab -e`)

```
# OPTIONAL:
# Daily job to create cached plots for the previous day to make loading
common plots
# quicker. Run a few hours after midnight so all data is uploaded.
0 4 * * * /usr/bin/hedgehogctl plotcache_generate_cached_plots -D >>
/var/log/hedgehog/plotcache_generate_cached_plots.log -D 2>&1
# Daily job to generate RSSAC reports. By default report is generated
# for a single day 1 week ago. Must be run after the
database_process_rssac_data script (see Data Manager cron jobs)
0 3 * * * /usr/bin/hedgehogctl rssac_generate_reports >>
/var/log/hedgehog/rssac_generate_reports.log 2>&1
```